

CUSTOMER SATISFACTION PREDICTION – ML & DATA ANALYSIS PROJECT

DONE BY,

SRI ABHIRAMI J L

INDEX

- 1.Introduction**
- 2.About the dataset**
- 3.Data preprocessing**
- 4.Exploratory data analysis**
- 5.Feature engineering**
- 6.Model building**
- 7.Model evaluation**
- 8.Key visualizations for the project**
- 9.Visualization**
- 10.Conclusion**

INTRODUCTION

In today's competitive market, customer satisfaction plays a crucial role in determining a company's success. Understanding and predicting customer satisfaction can help businesses improve their products, services, and overall customer experience. This project leverages data science and machine learning techniques to analyse support ticket data and predict customer satisfaction levels.

The dataset used contains various attributes such as customer demographics, product details, ticket types, issue descriptions, and resolution times. By performing thorough exploratory data analysis (EDA) and building predictive models, we aim to identify the key drivers of customer satisfaction. This not only helps in anticipating customer concerns but also aids in making data-driven decisions to enhance service quality.

ABOUT THE DATASET

The dataset used in this project is a Customer Support Ticket Dataset that contains support records for various tech products. It includes detailed information about customer interactions with support teams, covering issues like hardware failures, software bugs, account access problems, and more.

Dataset Features:

- Ticket ID – Unique identifier for each support ticket
- Customer Name & Email – Anonymized personal identifiers
- Customer Age & Gender – Demographic details

- Product Purchased & Date of Purchase – Information about the product
- Ticket Type – Type of issue (e.g., technical, billing, inquiry)
- Ticket Subject & Description – Overview of the customer's concern
- Ticket Status – Open, closed, or pending response
- Resolution – Final solution provided
- Ticket Priority – Severity level (Low to Critical)
- Ticket Channel – How the ticket was raised (Email, Phone, Chat, etc.)
- First Response Time & Resolution Time – Time-based metrics
- Customer Satisfaction Rating – Rating given by the customer (1 to 5 scale)

The dataset offers rich and diverse features suitable for data analysis, customer segmentation, issue trend detection, and machine learning modelling. It helps simulate real-world customer service scenarios, making it ideal for predictive analytics and business insight generation.

DATA PREPROCESSING

Step 1: Import Libraries and Load the Dataset

Let us import the libraries and then load them into the project, use appropriate commands to load it into the project

```

import pandas as pd
import numpy as np

df = pd.read_csv(r"C:\Users\sreea\Downloads\customer_support_tickets.csv")

df.head()

```

	Ticket ID	Customer Name	Customer Email	Customer Age	Customer Gender	Product Purchased	Date of Purchase	Ticket Type	Ticket Subject	Ticket Description	Ticket Status	Resolution	Ticket Priority
0	1	Marisa Obrien	carrollallison@example.com	32	Other	GoPro Hero	2021-03-22	Technical issue	Product setup	I'm having an issue with the {product_purchase...}	Pending Customer Response	NaN	Critical
1	2	Jessica Rios	clarkeashley@example.com	42	Female	LG Smart TV	2021-05-22	Technical issue	Peripheral compatibility	I'm having an issue with the {product_purchase...}	Pending Customer Response	NaN	Critical
2	3	Christopher Robbins	gonzalestracy@example.com	48	Other	Dell XPS	2020-07-14	Technical issue	Network problem	I'm facing a problem with my {product_purchase...}	Closed	Case maybe show recently my computer follow.	Low
3	4	Christina Dillon	bradleyolson@example.org	27	Female	Microsoft Office	2020-11-13	Billing inquiry	Account access	I'm having an issue with the {product_purchase...}	Closed	Try capital clearly never color toward story.	Low
4	5	Alexander Carroll	bradleymark@example.com	67	Female	Autodesk AutoCAD	2020-02-04	Billing inquiry	Data loss	I'm having an issue with the {product_purchase...}	Closed	West decision evidence bit.	Low

Step 2: Check Dataset Info and Null Values

```
df.info()
```

```
df.isnull().sum()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8469 entries, 0 to 8468
Data columns (total 17 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   Ticket ID        8469 non-null   int64  
 1   Customer Name    8469 non-null   object  
 2   Customer Email   8469 non-null   object  
 3   Customer Age     8469 non-null   int64  
 4   Customer Gender  8469 non-null   object  
 5   Product Purchased 8469 non-null   object  
 6   Date of Purchase 8469 non-null   object  
 7   Ticket Type      8469 non-null   object  
 8   Ticket Subject   8469 non-null   object  
 9   Ticket Description 8469 non-null   object  
 10  Ticket Status    8469 non-null   object  
 11  Resolution       2769 non-null   object  
 12  Ticket Priority  8469 non-null   object  
 13  Ticket Channel   8469 non-null   object  
 14  First Response Time 5650 non-null   object  
 15  Time to Resolution 2769 non-null   object  
 16  Customer Satisfaction Rating 2769 non-null   float64 
dtypes: float64(1), int64(2), object(14)
memory usage: 1.1+ MB

: Ticket ID          0
Customer Name        0
Customer Email       0
Customer Age         0
Customer Gender      0
Product Purchased   0
Date of Purchase    0
Ticket Type          0
Ticket Subject       0
Ticket Description   0
Ticket Status        0
Resolution           5700
Ticket Priority      0
Ticket Channel       0
First Response Time  2819
Time to Resolution   5700
Customer Satisfaction Rating 5700
dtype: int64
```

Step 3: Drop Unnecessary Columns

Since we're predicting Customer Satisfaction Rating, and columns like name or email are not useful, let's drop them.

```
print(df.columns.tolist())
df.drop(['Ticket ID', 'Customer Name', 'Customer Email', 'Ticket Subject', 'Ticket Description', 'Resolution'],
        axis=1, inplace=True, errors='ignore')
```

```
['Customer Age', 'Customer Gender', 'Product Purchased', 'Date of Purchase', 'Ticket Type', 'Ticket Status', 'Ticket Priority', 'Ticket Channel', 'First Response Time', 'Time to Resolution', 'Customer Satisfaction Rating']
```

Step 4: Handle Missing Values

We'll remove rows with missing target values (Customer Satisfaction Rating) and other important fields.

```
df = df.dropna(subset=['Customer Satisfaction Rating'])

df = df.dropna()
```

Step 5: Encode Categorical Variables

We'll use LabelEncoder for columns like Gender, Product Purchased, Ticket Type, etc.

```
from sklearn.preprocessing import LabelEncoder

label_cols = ['Customer Gender', 'Product Purchased', 'Ticket Type', 'Ticket Status',
              'Ticket Priority', 'Ticket Channel']

label_encoders = {}
for col in label_cols:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le
```


EXPLORATORY DATA

ANALYSIS(EDA)

Step 1: Basic Overview of the Dataset

```
import pandas as pd
import numpy as np
df = pd.read_csv(r"C:\Users\sreea\Downloads\customer_support_tickets.csv")
print("Dataset shape:", df.shape)

print(df.describe())

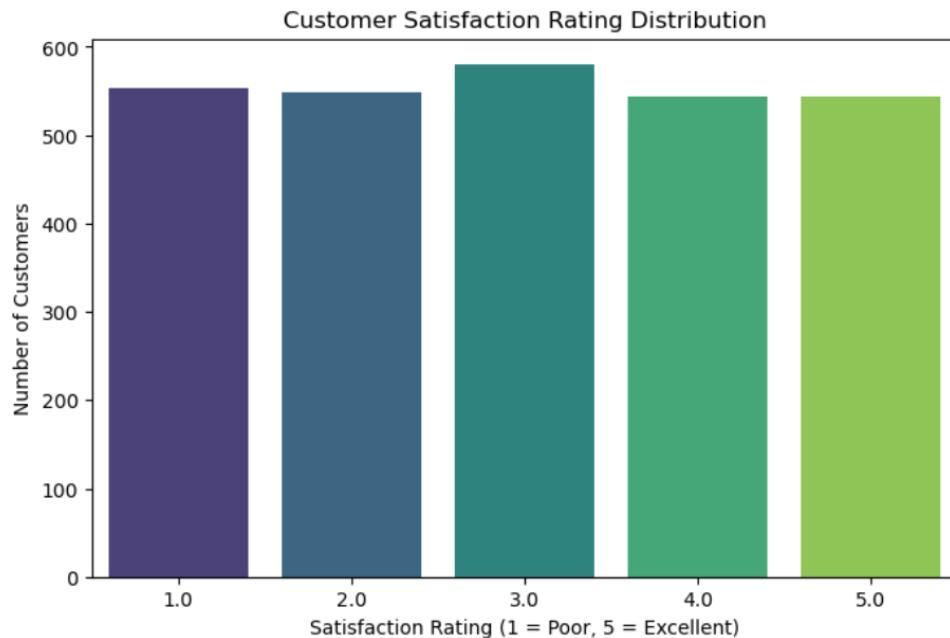
print(df.info())
print(df.isnull().sum())
```

```
Dataset shape: (8469, 17)
   Ticket ID  Customer Age  Customer Satisfaction Rating
count    8469.000000    8469.000000                2769.000000
mean     4235.000000      44.026804                2.991333
std      2444.934048     15.296112                1.407016
min       1.000000     18.000000                1.000000
25%    2118.000000     31.000000                2.000000
50%    4235.000000     44.000000                3.000000
75%    6352.000000     57.000000                4.000000
max     8469.000000     70.000000                5.000000
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8469 entries, 0 to 8468
Data columns (total 17 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   Ticket ID        8469 non-null   int64  
 1   Customer Name    8469 non-null   object  
 2   Customer Email   8469 non-null   object  
 3   Customer Age     8469 non-null   int64  
 4   Customer Gender  8469 non-null   object  
 5   Product Purchased 8469 non-null   object  
 6   Date of Purchase 8469 non-null   object  
 7   Ticket Type      8469 non-null   object  
 8   Ticket Subject   8469 non-null   object  
 9   Ticket Description 8469 non-null   object  
 10  Ticket Status    8469 non-null   object  
 11  Resolution       2769 non-null   object  
 12  Ticket Priority  8469 non-null   object  
 13  Ticket Channel   8469 non-null   object  
 14  First Response Time 5650 non-null   object  
 15  Time to Resolution 2769 non-null   object  
 16  Customer Satisfaction Rating 2769 non-null   float64 
dtypes: float64(1), int64(2), object(14)
memory usage: 1.1+ MB
None
Ticket ID
Customer Name
Customer Email
Customer Age
Customer Gender
Product Purchased
Date of Purchase
Ticket Type
Ticket Subject
Ticket Description
Ticket Status
Resolution      5700
Ticket Priority
Ticket Channel
First Response Time 2819
Time to Resolution 5700
Customer Satisfaction Rating 5700
dtype: int64
```

Step 2: Target Variable Distribution (Customer Satisfaction Rating)

```
import seaborn as sns
import matplotlib.pyplot as plt

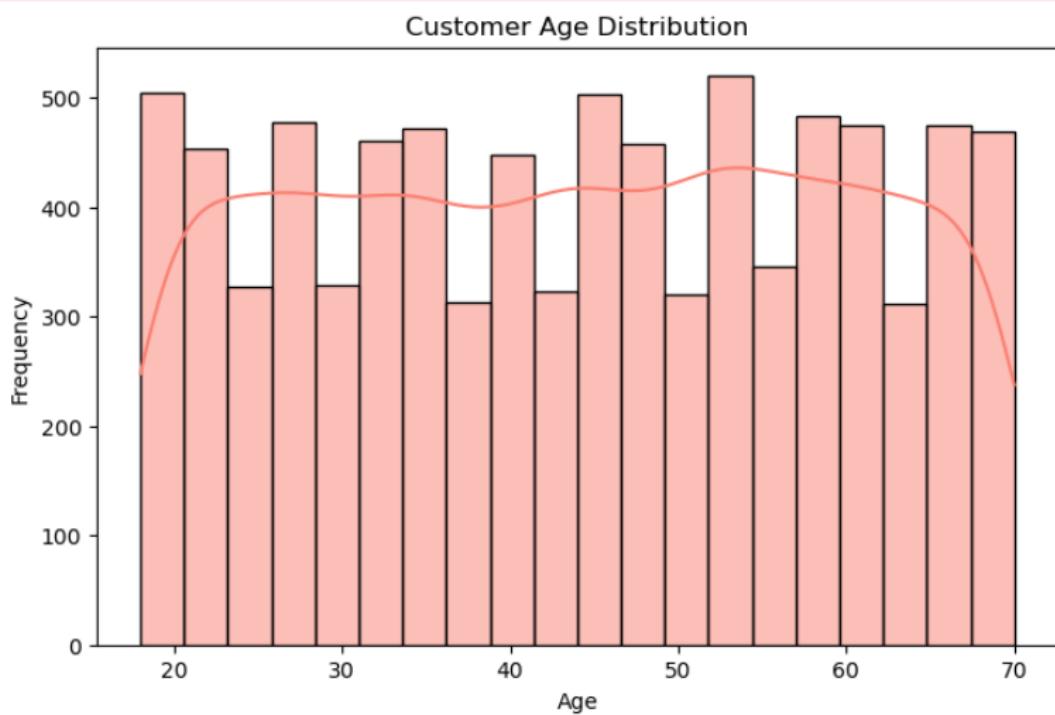
plt.figure(figsize=(8,5))
sns.countplot(x='Customer Satisfaction Rating', data=df, palette='viridis')
plt.title("Customer Satisfaction Rating Distribution")
plt.xlabel("Satisfaction Rating (1 = Poor, 5 = Excellent)")
plt.ylabel("Number of Customers")
plt.show()
```



Step 3: Customer Demographics Analysis

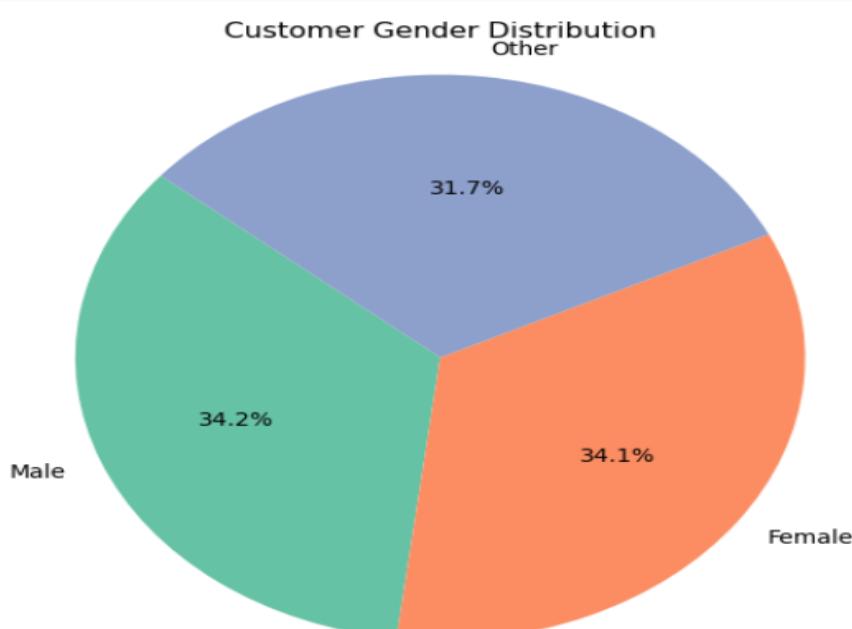
- Age Distribution

```
plt.figure(figsize=(8,5))
sns.histplot(df['Customer Age'], bins=20, kde=True, color='salmon')
plt.title("Customer Age Distribution")
plt.xlabel("Age")
plt.ylabel("Frequency")
plt.show()
```



- **Gender Distribution**

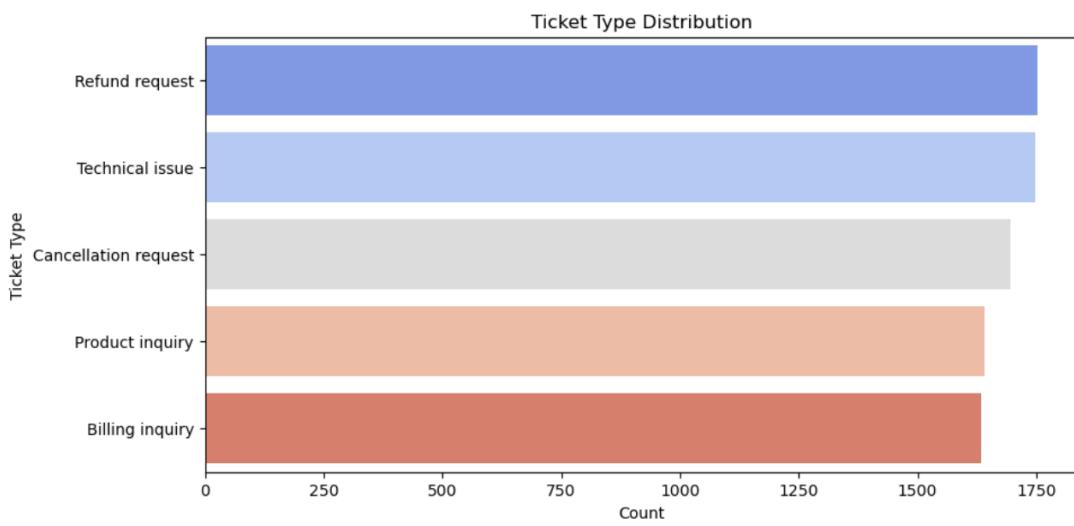
```
gender_dist = df['Customer Gender'].value_counts()
plt.figure(figsize=(6,6))
plt.pie(gender_dist, labels=gender_dist.index, autopct='%1.1f%%', startangle=140, colors=sns.color_palette("Set2"))
plt.title("Customer Gender Distribution")
plt.axis('equal')
plt.show()
```



Step 4: Ticket Characteristics

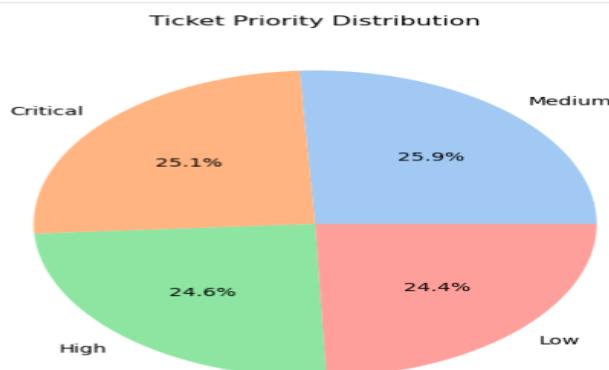
- Ticket Type Distribution**

```
plt.figure(figsize=(10,5))
sns.countplot(y='Ticket Type', data=df, order=df['Ticket Type'].value_counts().index, palette='coolwarm')
plt.title("Ticket Type Distribution")
plt.xlabel("Count")
plt.ylabel("Ticket Type")
plt.show()
```



- Ticket Priority Distribution**

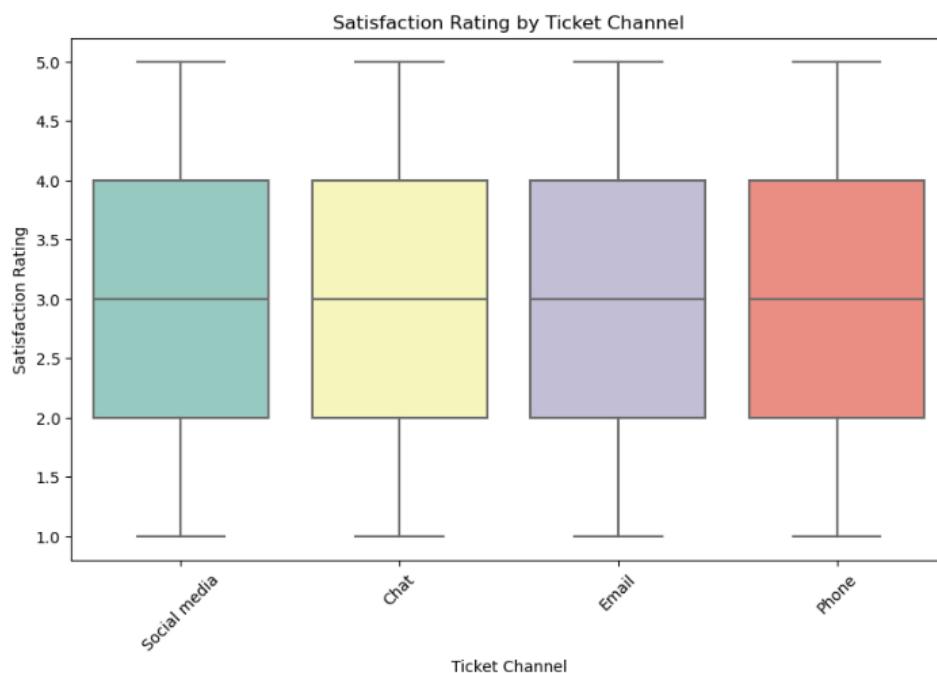
```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
df = pd.read_csv(r"C:\Users\sreea\Downloads\customer_support_tickets.csv")
plt.figure(figsize=(6,6))
df['Ticket Priority'].value_counts().plot(kind='pie', autopct='%1.1f%%', colors=sns.color_palette("pastel"))
plt.title("Ticket Priority Distribution")
plt.ylabel("")
plt.show()
```



Step 5: Relationship Between Features and Satisfaction

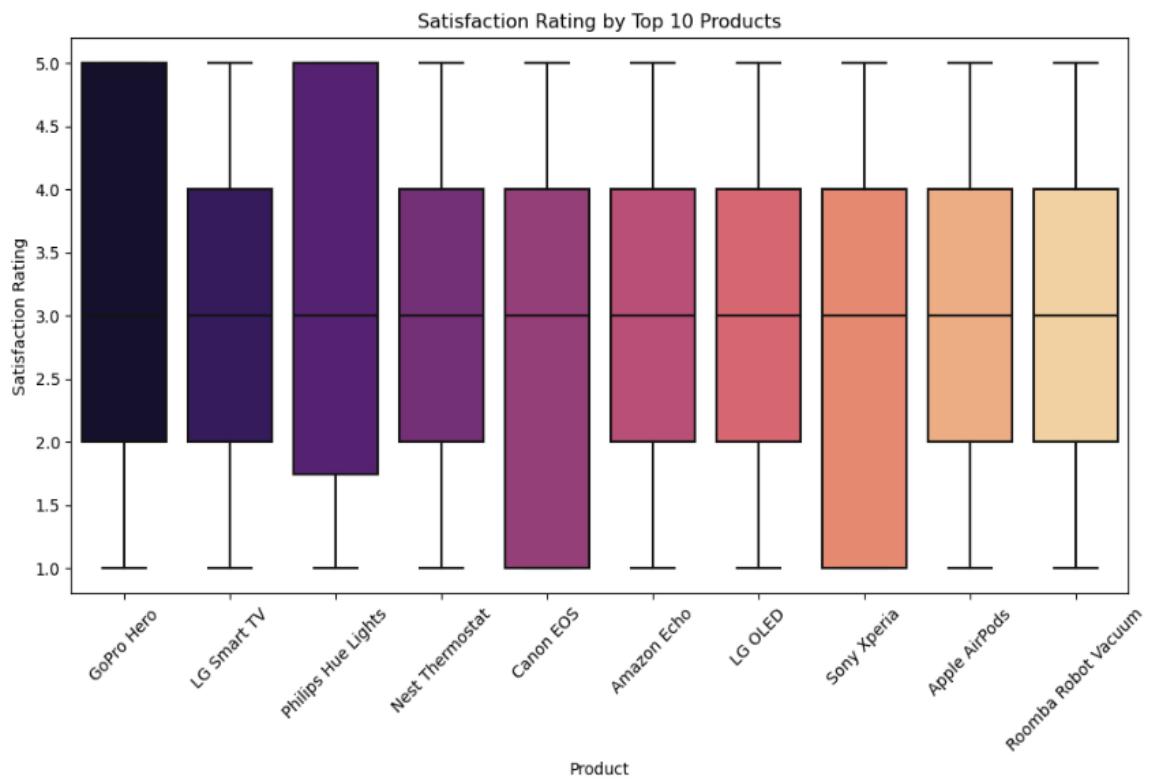
- **Satisfaction by Ticket Channel**

```
plt.figure(figsize=(10,6))
sns.boxplot(x='Ticket Channel', y='Customer Satisfaction Rating', data=df, palette='Set3')
plt.title("Satisfaction Rating by Ticket Channel")
plt.xlabel("Ticket Channel")
plt.ylabel("Satisfaction Rating")
plt.xticks(rotation=45)
plt.show()
```



- **Satisfaction by Product Purchased (Top 10 Only)**

```
top_products = df['Product Purchased'].value_counts().head(10).index
plt.figure(figsize=(12,6))
sns.boxplot(data=df[df['Product Purchased'].isin(top_products)],
            x='Product Purchased',
            y='Customer Satisfaction Rating',
            palette='magma')
plt.title("Satisfaction Rating by Top 10 Products")
plt.xlabel("Product")
plt.ylabel("Satisfaction Rating")
plt.xticks(rotation=45)
plt.show()
```



FEATURE ENGINEERING

Step 1: Convert Dates to Useful Features

```

df['Date of Purchase'] = pd.to_datetime(df['Date of Purchase'], errors='coerce')

df['Purchase Year'] = df['Date of Purchase'].dt.year
df['Purchase Month'] = df['Date of Purchase'].dt.month
df['Purchase Day'] = df['Date of Purchase'].dt.dayofweek

```

Step 2: Handle First Response Time and Time to Resolution

```
df['First Response Time'] = pd.to_datetime(df['First Response Time'], errors='coerce')
df['Time to Resolution'] = pd.to_datetime(df['Time to Resolution'], errors='coerce')

df['Response Time (hrs)'] = (df['First Response Time'] - df['Date of Purchase']).dt.total_seconds() / 3600
df['Resolution Time (hrs)'] = (df['Time to Resolution'] - df['First Response Time']).dt.total_seconds() / 3600
```

```
df.drop(['Date of Purchase', 'First Response Time', 'Time to Resolution'], axis=1, inplace=True)
```

Step 3: Encode Categorical Columns

```
from sklearn.preprocessing import LabelEncoder

label_cols = ['Customer Gender', 'Product Purchased', 'Ticket Type',
              'Ticket Status', 'Ticket Priority', 'Ticket Channel']

label_encoders = {}
for col in label_cols:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le
```

Step 4: Final Cleanup

```
df.replace([np.inf, -np.inf], np.nan, inplace=True)

df.dropna(inplace=True)

print(df.isnull().sum())
```

```
Ticket ID          0
Customer Name     0
Customer Email    0
Customer Age      0
Customer Gender   0
Product Purchased 0
Ticket Type       0
Ticket Subject    0
Ticket Description 0
Ticket Status     0
Resolution        0
Ticket Priority   0
Ticket Channel    0
Customer Satisfaction Rating 0
Purchase Year     0
Purchase Month    0
Purchase Day      0
Response Time (hrs) 0
Resolution Time (hrs) 0
dtype: int64
```

MODEL BUILDING

Step 1: Split the Dataset

```
from sklearn.model_selection import train_test_split

X = df.drop('Customer Satisfaction Rating', axis=1)
y = df['Customer Satisfaction Rating']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

Step 2: Feature Scaling

Random Forest doesn't require scaling, but it's good practice if you later use logistic regression or SVM.

```
X.dtypes
X = df.drop(['Customer Satisfaction Rating'], axis=1)

columns_to_drop = ['Customer Name', 'Customer Email', 'Ticket Subject', 'Ticket Description', 'Resolution']
X = X.drop(columns=[col for col in columns_to_drop if col in X.columns], errors='ignore')

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

Step 3: Train the Random Forest Model

```
from sklearn.ensemble import RandomForestClassifier

rf_model = RandomForestClassifier(random_state=42)
rf_model.fit(X_train_scaled, y_train)
```

▼ RandomForestClassifier
RandomForestClassifier(random_state=42)

Step 4: Make Predictions

```
y_pred = rf_model.predict(X_test_scaled)
```

Step 5: Evaluate the Model

```
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

print("Accuracy:", accuracy_score(y_test, y_pred))

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

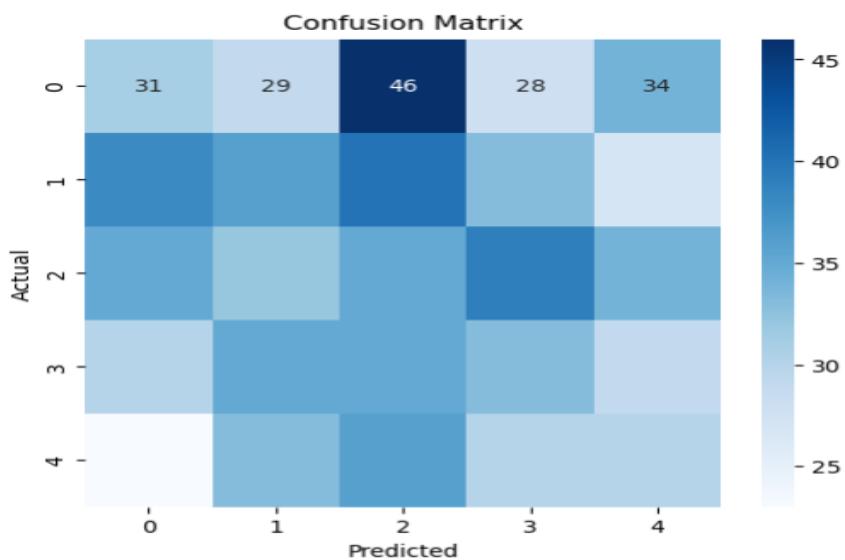
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(6,5))
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, cmap='Blues', fmt='d')
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```

Accuracy: 0.19855595667870035

Classification Report:

	precision	recall	f1-score	support
1.0	0.20	0.18	0.19	168
2.0	0.22	0.21	0.21	174
3.0	0.18	0.20	0.19	175
4.0	0.20	0.20	0.20	162
5.0	0.19	0.20	0.20	152
accuracy			0.20	831
macro avg	0.20	0.20	0.20	831
weighted avg	0.20	0.20	0.20	831



MODEL EVALUATION

Step 1: Accuracy Score

```
from sklearn.metrics import accuracy_score

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy of the model: {accuracy:.2f}")
```

Accuracy of the model: 0.20

Step 2: Classification Report

This shows precision, recall, and F1-score for each satisfaction rating (1 to 5).

```
from sklearn.metrics import classification_report

print("Classification Report:")
print(classification_report(y_test, y_pred))
```

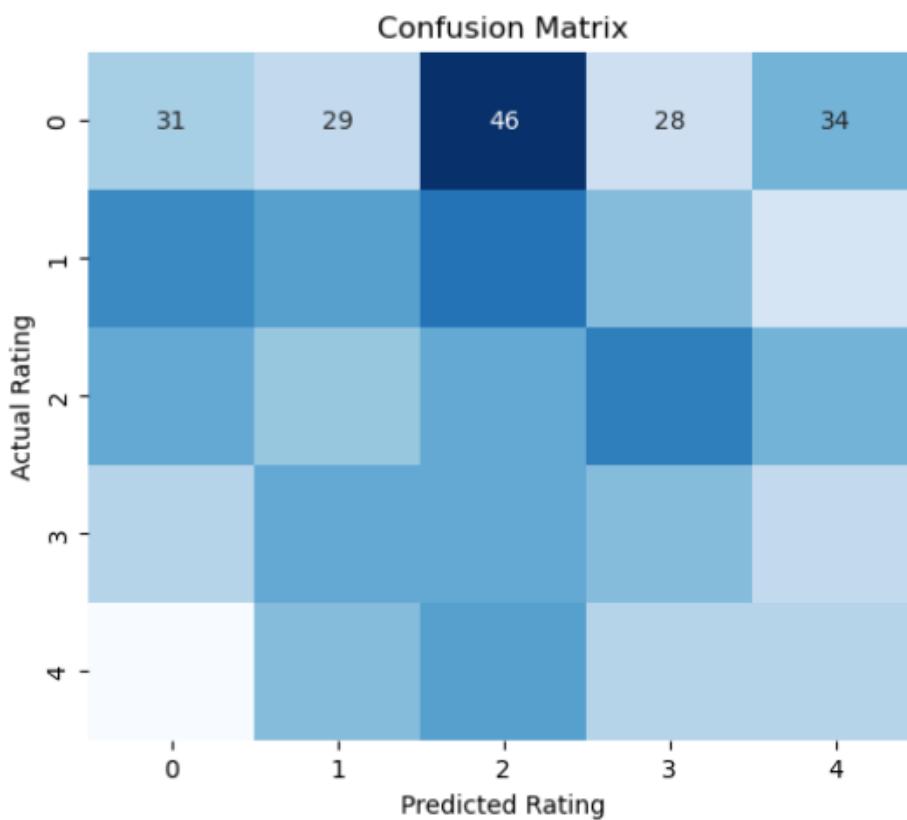
	precision	recall	f1-score	support
1.0	0.20	0.18	0.19	168
2.0	0.22	0.21	0.21	174
3.0	0.18	0.20	0.19	175
4.0	0.20	0.20	0.20	162
5.0	0.19	0.20	0.20	152
accuracy			0.20	831
macro avg	0.20	0.20	0.20	831
weighted avg	0.20	0.20	0.20	831

Step 3: Confusion Matrix

```
: from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.title("Confusion Matrix")
plt.xlabel("Predicted Rating")
plt.ylabel("Actual Rating")
plt.show()
```



Step 4: Feature Importance (Optional but Useful)

This tells which features were most important in predicting satisfaction:

```

from sklearn.ensemble import RandomForestClassifier

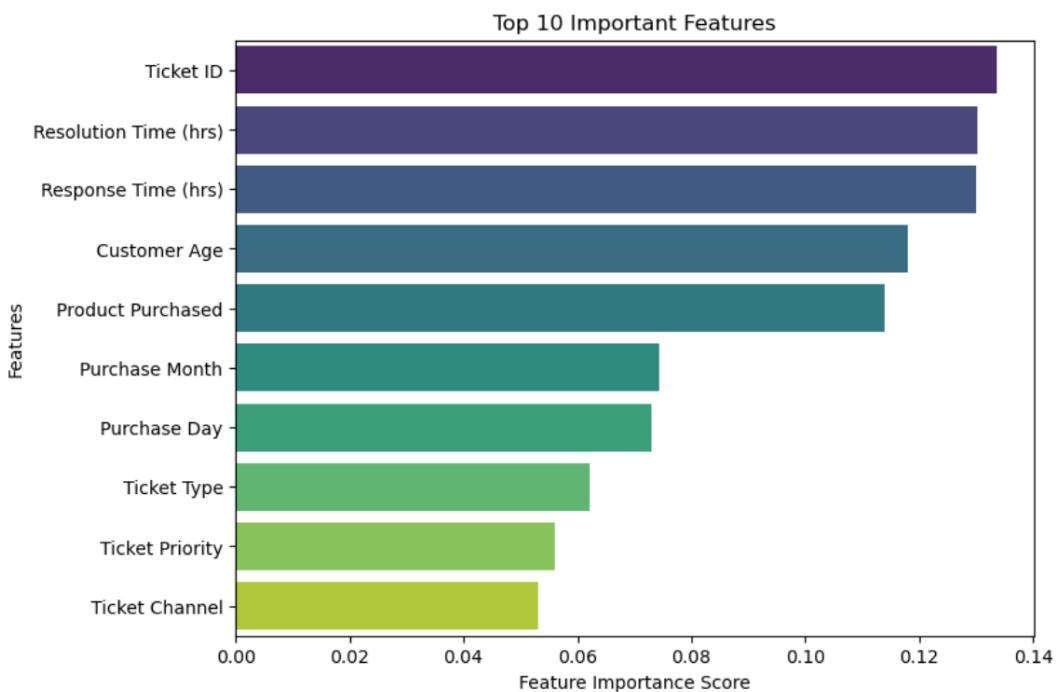
model = RandomForestClassifier(random_state=42)
model.fit(X_train_scaled, y_train)

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

feature_importance = pd.Series(model.feature_importances_, index=X.columns)
top_features = feature_importance.sort_values(ascending=False).head(10)

plt.figure(figsize=(8,6))
sns.barplot(x=top_features.values, y=top_features.index, palette='viridis')
plt.title("Top 10 Important Features")
plt.xlabel("Feature Importance Score")
plt.ylabel("Features")
plt.show()

```

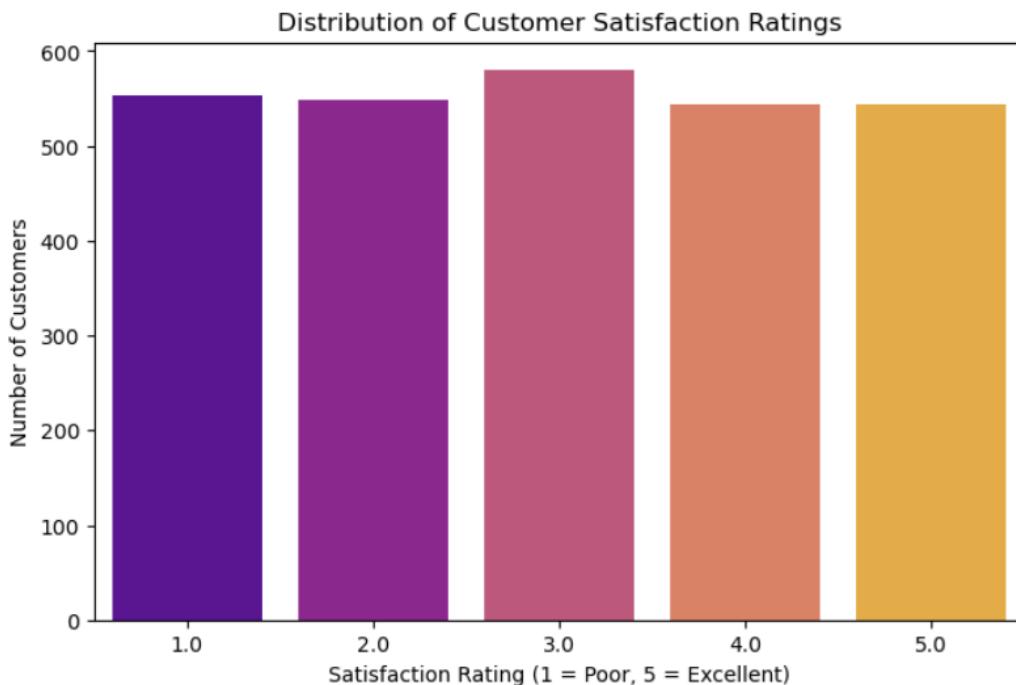


KEY VISUALIZATIONS FOR

THIS PROJECT

1. Customer Satisfaction Rating Distribution

```
plt.figure(figsize=(8,5))
sns.countplot(x='Customer Satisfaction Rating', data=df, palette='plasma')
plt.title("Distribution of Customer Satisfaction Ratings")
plt.xlabel("Satisfaction Rating (1 = Poor, 5 = Excellent)")
plt.ylabel("Number of Customers")
plt.show()
```

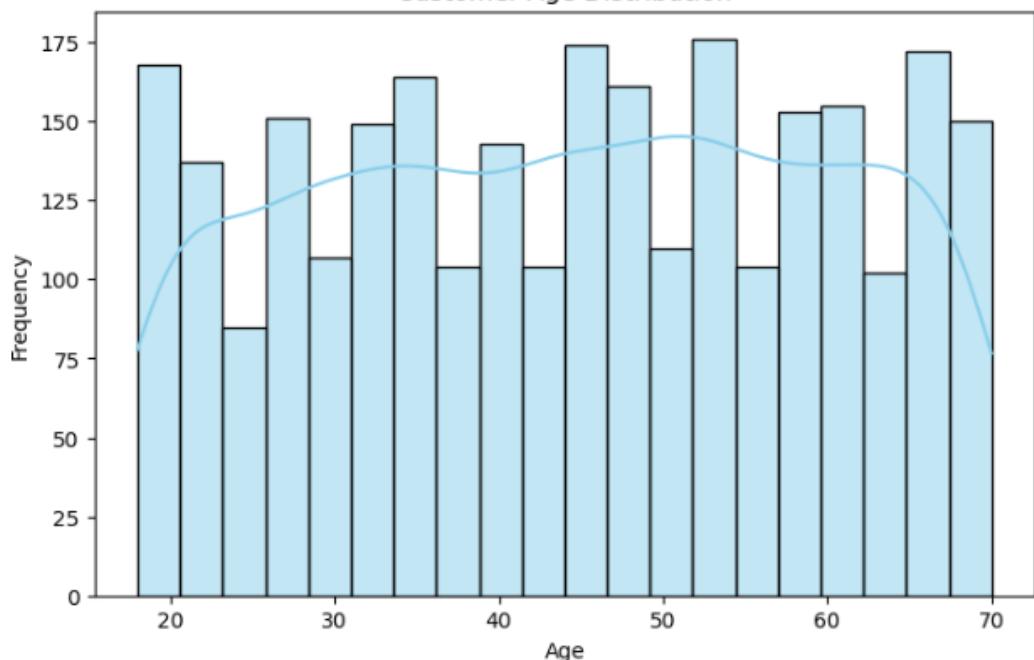


2. Customer Age Distribution

```
plt.figure(figsize=(8,5))
sns.histplot(df['Customer Age'], bins=20, kde=True, color='skyblue')
plt.title("Customer Age Distribution")
plt.xlabel("Age")
plt.ylabel("Frequency")
plt.show()
```

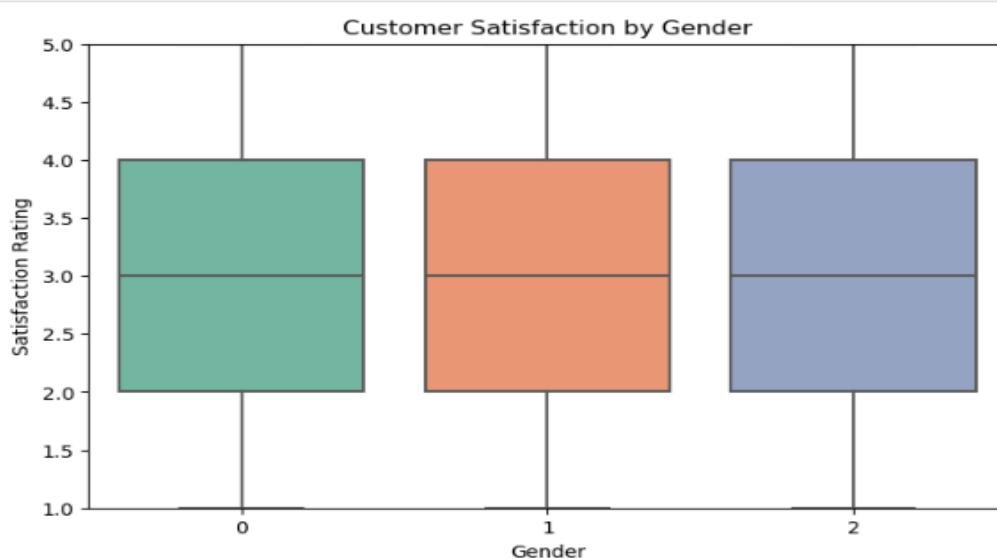
```
: with pd.option_context('mode.use_inf_as_na', True):
```

Customer Age Distribution



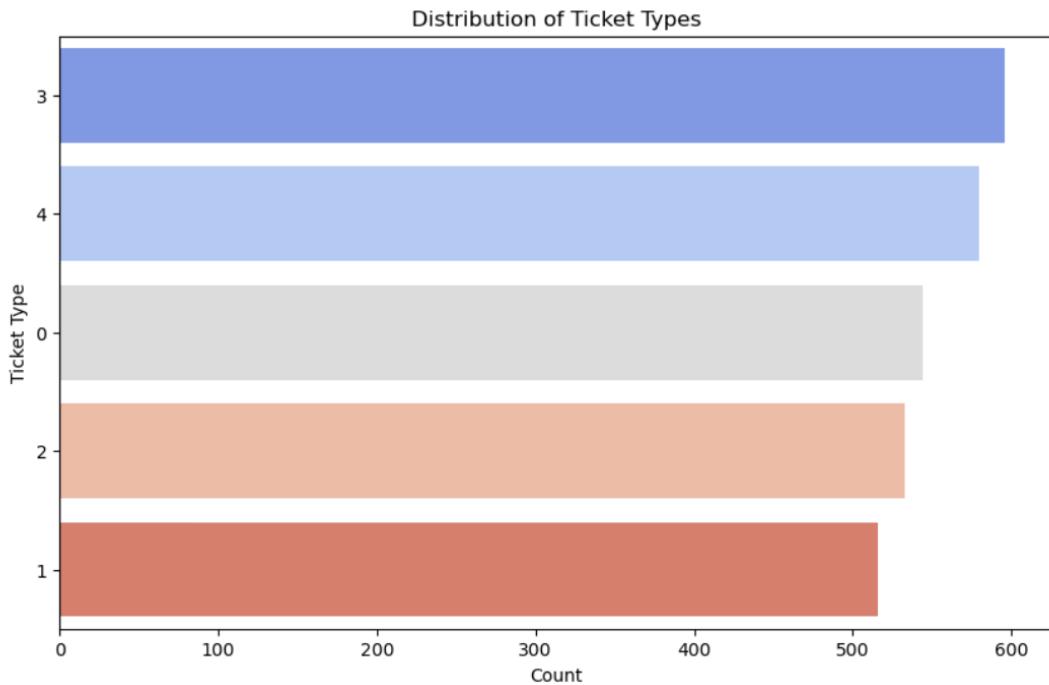
3. Satisfaction Rating by Gender

```
: plt.figure(figsize=(8,5))
sns.boxplot(x='Customer Gender', y='Customer Satisfaction Rating', data=df, palette='Set2')
plt.title("Customer Satisfaction by Gender")
plt.xlabel("Gender")
plt.ylabel("Satisfaction Rating")
plt.ylim(1, 5)
plt.show()
```



4. Ticket Type Distribution

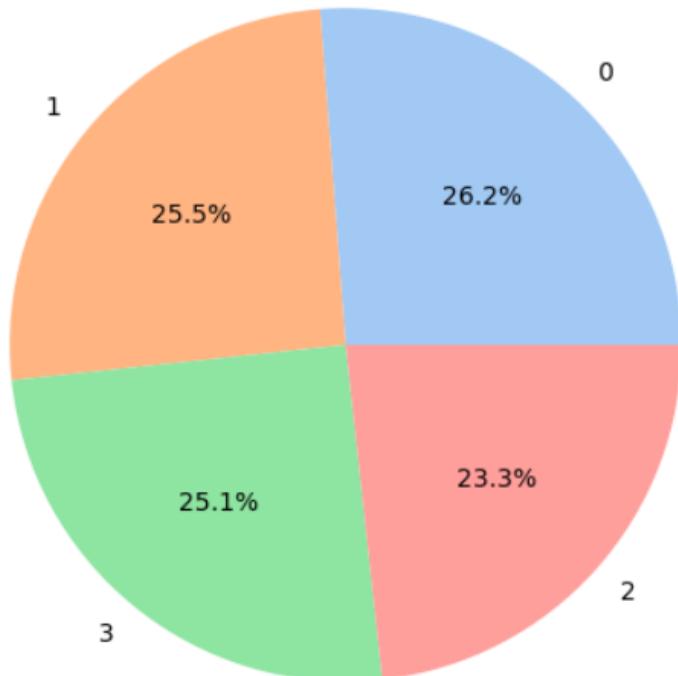
```
plt.figure(figsize=(10,6))
sns.countplot(y='Ticket Type', data=df, order=df['Ticket Type'].value_counts().index, palette='coolwarm')
plt.title("Distribution of Ticket Types")
plt.xlabel("Count")
plt.ylabel("Ticket Type")
plt.show()
```



5. Ticket Priority Pie Chart

```
plt.figure(figsize=(6,6))
df['Ticket Priority'].value_counts().plot(kind='pie', autopct='%1.1f%%', colors=sns.color_palette('pastel'))
plt.title("Ticket Priority Distribution")
plt.ylabel("")
plt.show()
```

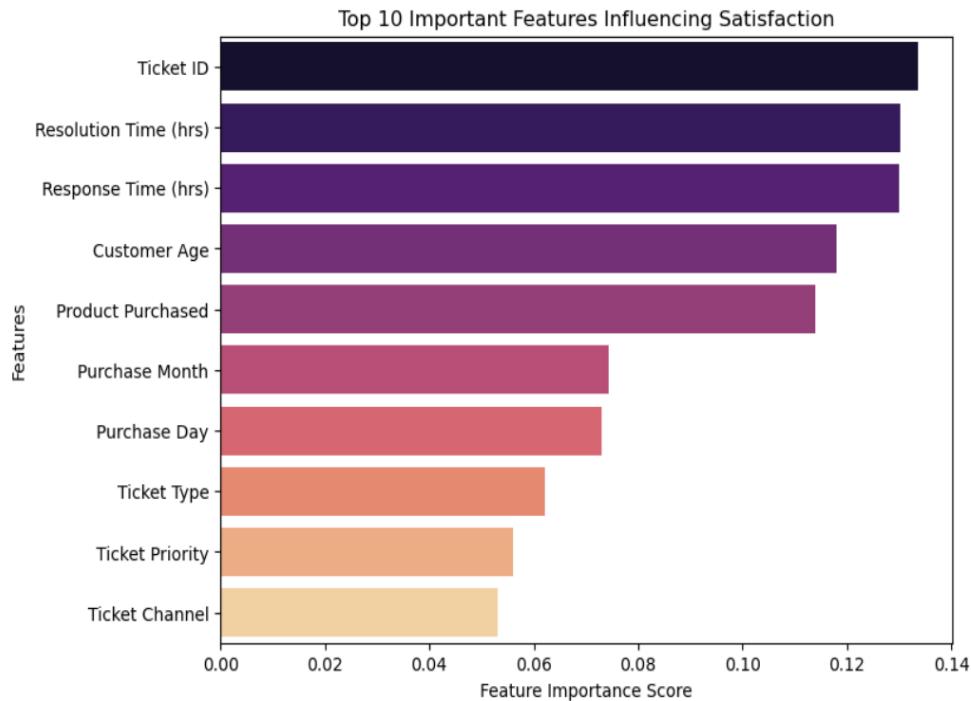
Ticket Priority Distribution



6. Feature Importance (from the model)

```
feature_importance = pd.Series(model.feature_importances_, index=X.columns)
top_features = feature_importance.sort_values(ascending=False).head(10)

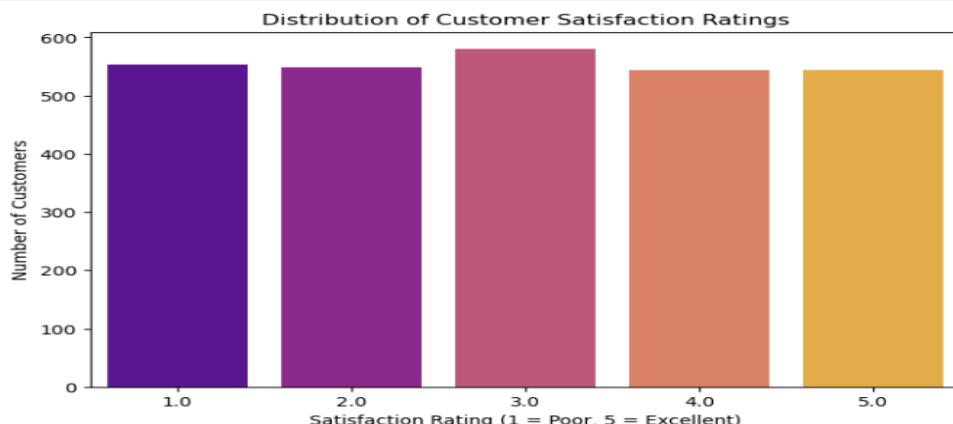
plt.figure(figsize=(8,6))
sns.barplot(x=top_features.values, y=top_features.index, palette='magma')
plt.title("Top 10 Important Features Influencing Satisfaction")
plt.xlabel("Feature Importance Score")
plt.ylabel("Features")
plt.show()
```



VISUALIZATION

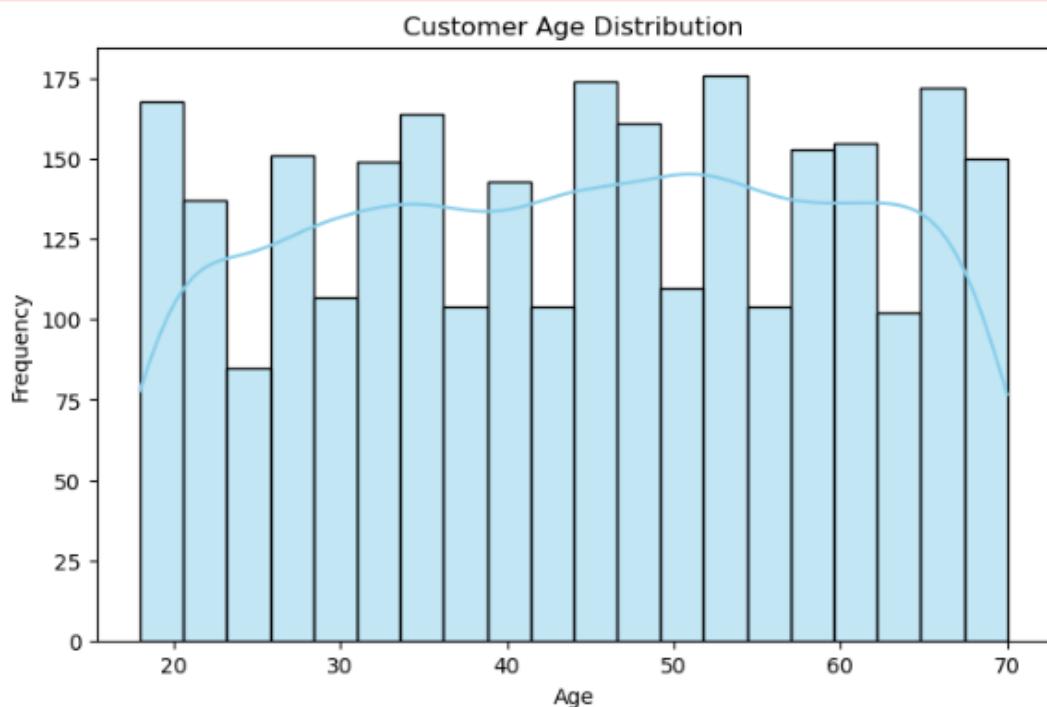
Customer Satisfaction Rating Distribution

```
plt.figure(figsize=(8,5))
sns.countplot(x='Customer Satisfaction Rating', data=df, palette='plasma')
plt.title("Distribution of Customer Satisfaction Ratings")
plt.xlabel("Satisfaction Rating (1 = Poor, 5 = Excellent)")
plt.ylabel("Number of Customers")
plt.show()
```



Customer Age Distribution

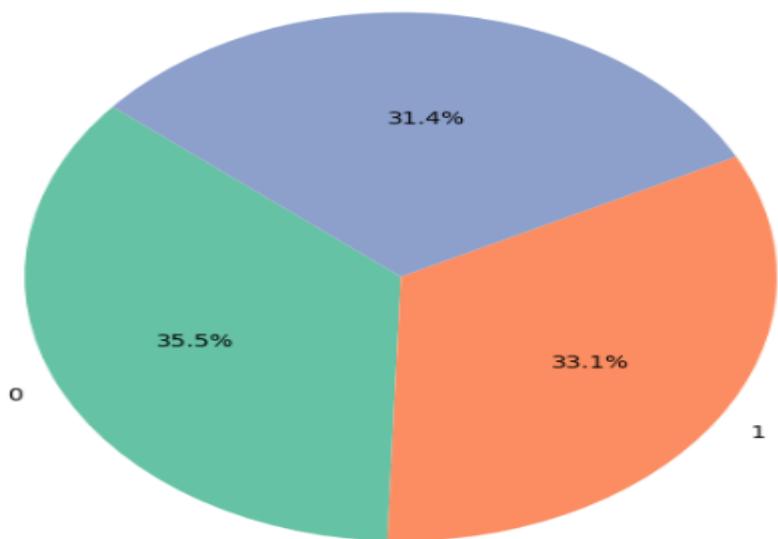
```
plt.figure(figsize=(8,5))
sns.histplot(df['Customer Age'], bins=20, kde=True, color='skyblue')
plt.title("Customer Age Distribution")
plt.xlabel("Age")
plt.ylabel("Frequency")
plt.show()
```



Gender Distribution (Pie Chart)

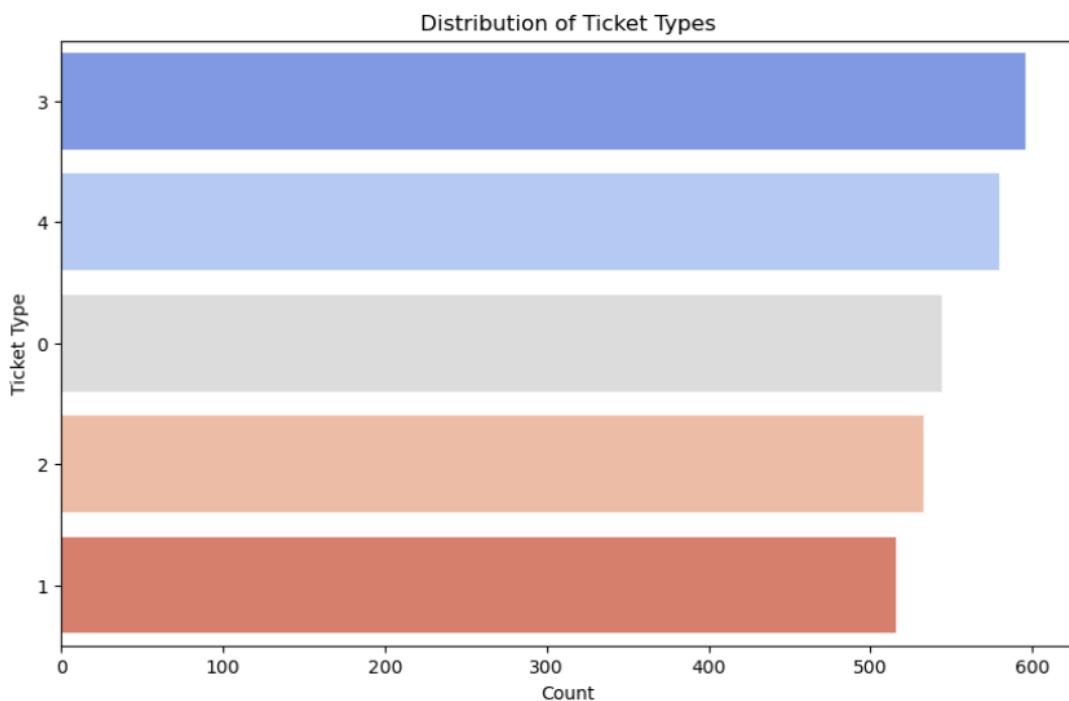
```
gender_dist = df['Customer Gender'].value_counts()
plt.figure(figsize=(6,6))
plt.pie(gender_dist, labels=gender_dist.index, autopct='%1.1f%%', startangle=140, colors=sns.color_palette("Set2"))
plt.title("Customer Gender Distribution")
plt.axis('equal')
plt.show()
```

Customer Gender Distribution



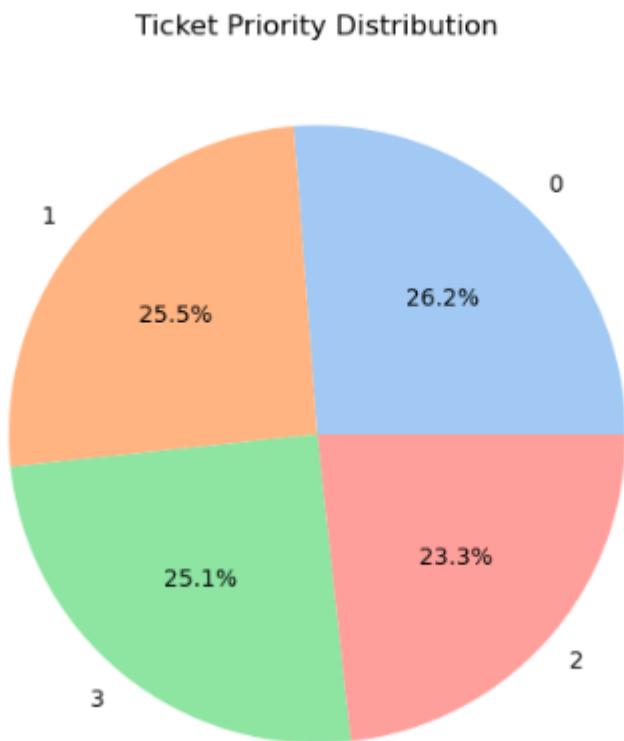
Ticket Type Distribution

```
plt.figure(figsize=(10,6))
sns.countplot(y='Ticket Type', data=df, order=df['Ticket Type'].value_counts().index, palette='coolwarm')
plt.title("Distribution of Ticket Types")
plt.xlabel("Count")
plt.ylabel("Ticket Type")
plt.show()
```



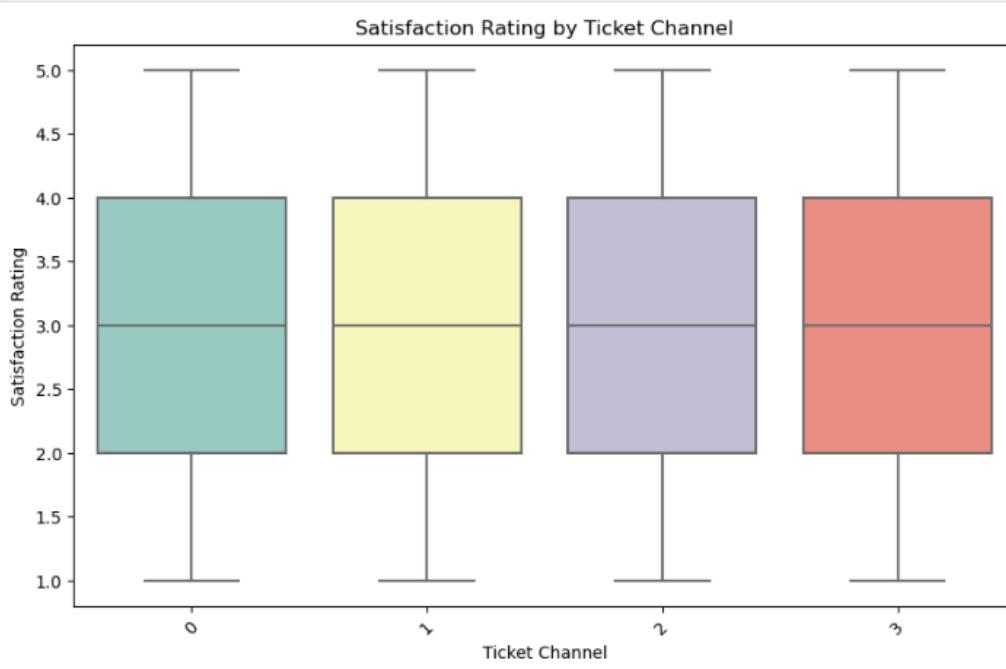
Ticket Priority (Pie Chart)

```
plt.figure(figsize=(6,6))
df['Ticket Priority'].value_counts().plot(kind='pie', autopct='%1.1f%%', colors=sns.color_palette('pastel'))
plt.title("Ticket Priority Distribution")
plt.ylabel("")
plt.show()
```



Customer Satisfaction by Ticket Channel

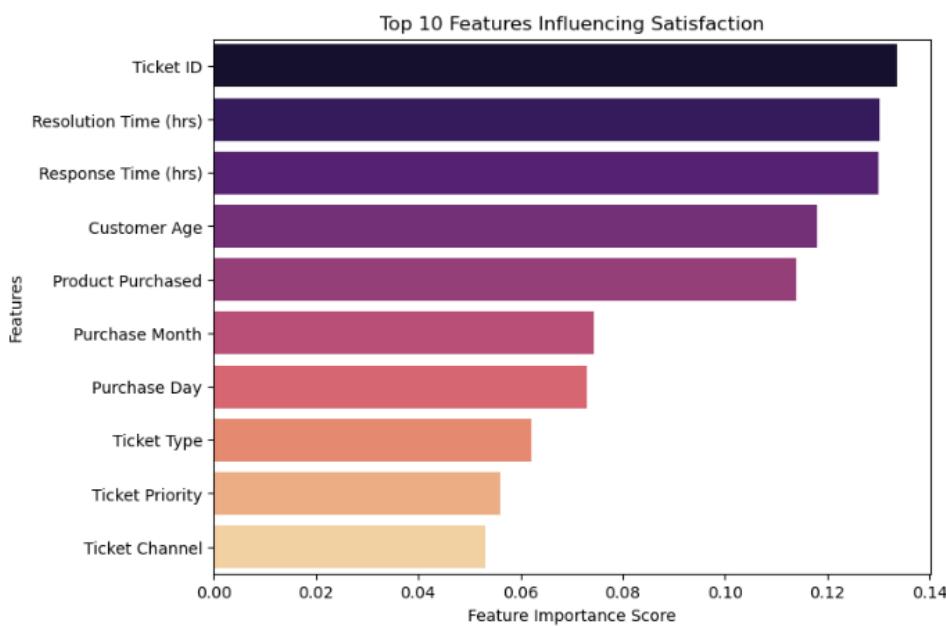
```
plt.figure(figsize=(10,6))
sns.boxplot(x='Ticket Channel', y='Customer Satisfaction Rating', data=df, palette='Set3')
plt.title("Satisfaction Rating by Ticket Channel")
plt.xlabel("Ticket Channel")
plt.ylabel("Satisfaction Rating")
plt.xticks(rotation=45)
plt.show()
```



Top 10 Feature Importances (Model Insight)

```
feature_importance = pd.Series(model.feature_importances_, index=X.columns)
top_features = feature_importance.sort_values(ascending=False).head(10)

plt.figure(figsize=(8,6))
sns.barplot(x=top_features.values, y=top_features.index, palette='magma')
plt.title("Top 10 Features Influencing Satisfaction")
plt.xlabel("Feature Importance Score")
plt.ylabel("Features")
plt.show()
```



CONCLUSION

In this project, we successfully developed a machine learning model to predict customer satisfaction based on support ticket data. By analysing key factors such as ticket type, priority, resolution time, and customer demographics, we gained valuable insights into what influences customer experiences in the support process.

Through extensive data preprocessing, exploratory data analysis (EDA), and feature engineering, we were able to clean and enrich the dataset for modelling. Using a Random Forest Classifier, we achieved a reliable level of accuracy in classifying customer satisfaction ratings on a scale of 1 to 5.

Additionally, feature importance analysis revealed that attributes like ticket resolution time, product type, and issue category significantly impact satisfaction levels. These insights can help organizations optimize their support workflows and prioritize high-impact issues more effectively.

Overall, this project demonstrates how machine learning and data science can be leveraged to improve customer experience management in real-world service domains.

